
DEFIANCE User Documentation

Release 1.0

BPHK2023

Aug 16, 2024

CONTENTS

1	Overview	1
2	RL-Applications	3
2.1	RLApplication	3
2.2	AgentApplication	3
2.3	ObservationApplication	7
2.4	RewardApplication	9
2.5	ActionApplication	9
2.6	Communication between RL-Applications	10
3	Data History Container	13
4	ChannelInterface	15
4.1	Overview	15
4.2	Usage	16
4.3	SimpleChannelInterface	16
4.4	SocketChannelInterface	17
4.5	Custom Channel Interface	19
5	Helper Classes	21
5.1	RLApplicationHelper	21
5.2	CommunicationHelper	22

OVERVIEW

The *ns-3* DEFIANCE module is a reinforcement learning (RL) framework for *ns-3*. It allows the simulation of distributed RL in networks. It can handle single-agent and multi-agent RL scenarios.

The user performs the following steps to carry out the RL experiment:

1. Implement the network topology and traffic using standard *ns-3* code.
2. To define how observations and reward signals are collected, how actions are executed in the environment and how the agents perform inference and training, subclass from the provided *RL-Applications*. Abstract classes for the different subtasks are provided via the *AgentApplication*, *ObservationApplication*, *RewardApplication* and *ActionApplication*. These applications are installed in the simulation via the *RLApplicationHelper*.
3. Specify how data is exchanged between these components and specify the communication structure via channels. The lowest level of abstraction our framework proposes is using *ChannelInterfaces* for this. The framework also provides a *CommunicationHelper* class to simplify the communication setup.
4. Finally, use the utilities provided by *ns3-ai* to interact with the simulation as an RL environment.

Note: An in-depth documentation of the multi-agent interface we added to *ns3-ai* can be found [here](#). This interface is also used by the DEFIANCE framework. In case you are interested in how the framework functions or think about extending DEFIANCE, we recommend to take a look at these docs or the *blog post* <https://medium.com/@oliver.zimmermann/reinforcement-learning-in-ns3-part-1-698b9c30c0cd> we wrote about it.

RL-APPLICATIONS

2.1 RlApplication

All presented applications inherit from `RlApplication`. Therefore, all of them can access the following functionality.

2.1.1 RlApplicationId

A common use case is to identify the `RlApplication` information has been sent to or received from. To accomplish this, each `RlApplication` has a unique identifier, their `RlApplicationId`. This ID can be set by passing an `RlApplicationId` to `RlApplication::SetId`. `RlApplicationId` is a struct consisting of an `uint32_t` `applicationId` and an `ApplicationType` `applicationType`.

This introduces four different types of `RlApplication`:

- OBSERVATION
- REWARD
- AGENT
- ACTION

If the provided `CommunicationHelper` is not used, one must set all `RlApplicationIds` manually to be able to use inter-application communication.

2.2 AgentApplication

The `AgentApplication` is where inference is performed. This code is an example of a possible implementation followed by an explanation.

```
1 class InferenceAgentApp : public AgentApplication
2 {
3 public:
4     InferenceAgentApp()
5         : AgentApplication()
6     {
7     };
8
9     ~InferenceAgentApp() override{};
10
11     static TypeId GetTypeId()
12     {
```

(continues on next page)

(continued from previous page)

```

13     static TypeId tid = TypeId("ns3::InferenceAgentApp")
14         .SetParent<AgentApplication>()
15         .SetGroupName("defiance")
16         .AddConstructor<InferenceAgentApp>();
17     return tid;
18 }
19
20 void Setup() override
21 {
22     AgentApplication::Setup();
23     m_observation = GetResetObservation();
24     m_reward = GetResetReward();
25 }
26
27 void OnRecvObs(uint id) override
28 {
29     m_observation = m_obsDataStruct.GetNewestByID(id)
30         ->data->Get("floatObs")
31         ->GetObject<OpenGymBoxContainer<float>>();
32     InferAction();
33 }
34
35 void OnRecvReward(uint id) override
36 {
37     m_reward = m_rewardDataStruct.GetNewestByID(0)
38         ->data->Get("reward")
39         ->GetObject<OpenGymBoxContainer<float>>()
40         ->GetValue(0);
41 }
42
43 Ptr<OpenGymDataContainer> GetResetObservation()
44 {
45     // This method returns the initial observation that is used after resetting
46     ↪the environment.
47     uint32_t shape = 4;
48     std::vector<uint32_t> vShape = {shape};
49     auto obj = CreateObject<OpenGymBoxContainer<float>>(vShape);
50     for(auto i = 0; i < shape; i++) {obj->AddValue(0);}
51     return obj;
52 }
53
54 float GetResetReward()
55 {
56     // This method returns the initial reward that is used after resetting the
57     ↪environment.
58     return 0.0;
59 }
60
61 private:
62 Ptr<OpenGymSpace> GetObservationSpace() override
63 {
64     uint32_t shape = 4;
65     std::vector<uint32_t> vShape = {shape};
66     std::string dtype = TypeNameGet<float>();
67
68     std::vector<float> low = {-4.8 * 2, -INFINITY, -0.418 * 2, -INFINITY};
69     std::vector<float> high = {4.8 * 2, INFINITY, 0.418 * 2, INFINITY};

```

(continues on next page)

(continued from previous page)

```

68         return CreateObject<OpenGymBoxSpace>(low, high, vShape, dtype);
69     }
70
71     Ptr<OpenGymSpace> GetActionSpace() override
72     {
73         return MakeBoxSpace<int>(1, 0, 1);
74     }
75 }
76 };

```

To implement your own `AgentApplication` it is necessary to inherit from `AgentApplication` in order to access all features provided by our framework. This can be seen in line 1. The method `InferenceAgentApp::GetTypeId` (lines 11-18) is mandatory, as it is part of the *ns-3* library. Since our classes inherit from `ns3::Object` one has to provide this method to allow the usage of *ns-3*-factories and *ns-3*-pointers.

`InferenceAgentApp::Setup` is called at the beginning of the scenario and ensures that all required variables for inference are initialized. It is advised to call the parent method (line 22) since it informs the MARL interface about the action and observation-space provided by `InferenceAgentApp::GetObservationSpace` and `InferenceAgentApp::GetActionSpace`. Additionally the `InferenceAgentApp::Setup` method can be used to initialize `m_observation` and `m_reward` since this method should always be called before the first occurrence of inference and thereby guarantees that no uninitialized variables will be used for inference.

`m_observation` and `m_reward` are two inherited variables from the `AgentApplication` class. `m_observation` is an `OpenGymDataContainer` that stores the observations used for inference. `m_reward` is simply a float value representing the current reward. Both of this variables are passed to the MARL interface when `AgentApplication::InferAction` is called (line 32).

`InferenceAgentApp::OnRecvObs` and `InferenceAgentApp::OnRecvReward` are called when the `AgentApplication` receives an observation or reward. The `id` is the ID of the `RLApplication` that sent the data. It can be used to retrieve the desired data from `m_obsDataStruct` or `m_rewardDataStruct` by calling `HistoryContainer::GetNewestByID(id)` (line 29). However, there is no restriction on how to update `m_observation` or whether `InferenceAgentApp::OnRecvObs` should be used at all.

Both of the data structures `m_obsDataStruct` and `m_rewardDataStruct` are instances of type `HistoryContainer`. Once a reward or observation is received, the `AgentApplication` ensures both are updated accordingly before calling `InferenceAgentApp::OnRecvObs` or `InferenceAgentApp::OnRecvReward`.

In line 32, the method `AgentApplication::InferAction` is called. As mentioned earlier passes this method all required parameters for inference to the MARL interface. Additionally, a callback is passed on that sends the returned action from the Python side to an `ActionApplication`. Pass the `RLApplicationId::applicationId` to `AgentApplication::InferAction` as in `InferAction(id)` if the received action should only be sent to a specific `ActionApplication`. Otherwise the action will be sent to all registered instances. It is not required to call this method in `InferenceAgentApp::OnRecvObs`. For example `AgentApplication::InferAction` could also be called in a method that is scheduled at equally spaced timesteps or after an *ns-3*-event. If preferred, it is even possible to call inference outside of the `ActionApplication`. Since `AgentApplication::InferAction` is by design protected within `AgentApplication`, this would require the usage of `OpenGymMultiAgentInterface::NotifyCurrentState` and thus thorough testing.

`InferenceAgentApp::OnRecvReward` is similar to `InferenceAgentApp::OnRecvObs` in terms of when it is called and its purpose. Both of these methods allow to aggregate over the data received by multiple `RLApplication` instances. One example could calculate the min of all rewards sent by `RewardApplication` in the method `InferenceAgentApp::OnRecvReward`.

`InferenceAgentApp::GetResetObservation` and `InferenceAgentApp::GetResetReward` are vital after a reset of the environment. Therefore, they must be implemented in the inheriting class. When setting up a scenario without training a Ray agent and no resets, they are optional, yet it is essential to initialize `m_observation` and `m_reward` at the beginning of the scenario (e.g. ll.23-24).

The last two important methods are `InferenceAgentApp::GetObservationSpace` and `InferenceAgentApp::GetActionSpace`. These methods are mandatory because they inform the MARL interface about the dimensions of the respective spaces. Information about the different spaces have to be provided in instances of `OpenGymSpace`. An exemplary creation of such spaces can be seen in line 62 to 69. These spaces as well as the `OpenGymDataContainer` are part of the *ns-3-ai* library. To reduce the overhead of creating an `OpenGymSpace` or `OpenGymDataContainer`, some useful functions are provided in `base-test.h`. An example usage of one of these functions can be seen in line 74.

2.2.1 Additional Features and Use-Cases

Configure History Containers

The length of `m_rewardDataStruct` and `m_obsDataStruct` can be changed by setting the attribute `MaxRewardHistoryLength` or `MaxObservationHistoryLength`.

It is also possible to save a timestamp, marking the time of arrival in `m_rewardDataStruct` and `m_obsDataStruct`. If this feature is required, set `ObservationTimestamping` or `RewardTimestamping` to `true`. More information is given in [Data History Container](#).

Provide Extra Info

To pass extra info to the environment, override the method `AgentApplication::GetExtraInfo`

```
/* ... */
private:
    std::string m_importantMessage;
    std::map<std::string, std::string> GetExtraInfo() override
    {
        std::map<std::string, std::string> info;
        info["agent"] = m_importantMessage;
        return info;
    }
```

Action Delay

To simulate the time required to calculate inference, a delay can be set between receiving an action and performing the callback specified for action execution.

```
/* ... */
private:
    std::string m_importantMessage;
    Time GetActionDelay() override
    {
        return Seconds(1);
    }
```

Override initiateAction and initiateActionForApp

After inference took place, either of these methods is invoked with the returned action from the MARL interface. This method then sends the received message to either all registered `ActionApplication` or the one that matches `remoteAppId`. Overriding this method allows for example to only send over a specific `ChannelInterface`.

```
/* in your AgentApplication-class: */
protected:
    uint32_t interfaceToUse;
    void InitiateActionForApp(uint remoteAppId, Ptr<OpenGymDataContainer> action)
    {
        SendAction(MakeDictContainer("default", action), remoteAppId, interfaceToUse);
    }
```

OnRecvFromAgent

To specify how an `AgentApplication` should handle messages from another `AgentApplication`, override this method.

The method receives a `remoteAppId` matching the `RlApplicationId::applicationId` of the `AgentApplication` that send the data and the message itself as a `Ptr<OpenGymDictContainer>`. Here is an example for this:

```
/* in your AgentApplication-class: */
protected:
    uint32_t agentOfInterest;
    void OnRecvFromAgent(uint remoteAppId, Ptr<OpenGymDictContainer> payload)
    {
        if(remoteAppId == agentOfInterest)
        {
            message = payload->Get("parameter")
                ->GetObject<OpenGymBoxContainer<float>>()
                ->GetValue(0);
        }
    }
```

If desired, a new `HistoryContainer` can be added to the class which can be used to store and retrieve the received agent messages in a similar fashion as the observations and rewards.

2.3 ObservationApplication

The main purpose of the `ObservationApplication` is to send observations to the agent. Therefore, the class is equipped with the methods `ObservationApplication::RegisterCallbacks` and `ObservationApplication::Send`. To implement an `ObservationApplication`, create a child class that inherits from `ObservationApplication`. This also requires overriding `GetTypeId` in a similar fashion as seen earlier in the `AgentApplication` example.

2.3.1 ObservationApplication::RegisterCallbacks

This method allows registration of callbacks to trace sources. This ensures the `ObservationApplication` is always informed when a value that should be observed changes.

```
class YourImplementation : public ObservationApplication{
public:
    /* ... */
    void
    RegisterCallbacks() override
    {
        DynamicCast<YourNode>(GetNode())->m_reportYourTrace.ConnectWithoutContext(
            MakeCallback(&YourImplementation::Observe, this));
    }
    void Observe(/*values provided by the trace source*/)
    {
        /* send observation or wait for more observation */
    }
}
```

Note: It can be tricky to access the required trace source inside the `ObservationApplication` class, especially if the trace source is not provided by *ns-3*. In this example, the costume trace source is accessed by inheriting the `Node` class and adding the trace source as a class member. All `ns3::Application` instances can access the node they are installed on with `GetNode`. Alternatively, trace sources can be accessed by a *ns-3* path. Look into the *ns-3* documentation for more information.

2.3.2 ObservationApplication::Send

Once the `ObservationApplication` is satisfied with the observations, it can send these observations to registered instances of `AgentApplication`. This functionality is offered by the base class. The observations have to be wrapped into an `OpenGymDictContainer`. If an observation should only be sent to a specific agent, pass the `RlApplicationId` to `ObservationApplication::Send`. Furthermore, the ID of the `ChannelInterface` can be provided. If not provided, the observation is sent to all registered instances.

```
class YourImplementation : public ObservationApplication{
public:
    /* ... */
    void Observe(uint32_t value)
    {
        /*create OpenGymDataContainer */

        Send(/*OpenGymDictContainer*/);
        // or
        Send(/*OpenGymDictContainer*/, remoteId, interfaceId);
    }
}
```

2.4 RewardApplication

The `RewardApplication` is in its functionality similar to `ObservationApplication` since both classes inherit from the same base class. A reward should be sent to an instance of `AgentApplication` once a relevant event is triggered. To accomplish that the `RewardApplication::Send` is provided. It is required to wrap all reward information into an `OpenGymDictContainer`.

```
class YourImplementation : public RewardApplication{
public:
    /* ... */
    void
    RegisterCallbacks() override
    {
        DynamicCast<YourNode>(GetNode())->m_reportYourTrace.ConnectWithoutContext(
            MakeCallback(&YourImplementation::ObserveReward, this));
    }
    void ObserveReward(/*values provided by the trace source*/)
    {
        /*create OpenGymDataContainer */

        Send(/*OpenGymDictContainer*/);
        // or
        Send(/*OpenGymDictContainer*/, remoteID, interfaceId);
    }
}
```

2.5 ActionApplication

The `ActionApplication` receives actions and executes them. Therefore, upon receiving an action from an `AgentApplication`, the virtual method `ActionApplication::ExecuteAction` is triggered. To specify what action should be performed, override the `ActionApplication::ExecuteAction` in a child class.

2.5.1 ActionApplication::ExecuteAction

In this method, two parameters are accessible. `remoteAppId` is the `RLApplicationId` of the `AgentApplication` that sent the action. `action` is an `OpenGymDictContainer` that contains the sent action out of the action space. An exemplary retrieval of the actual content of action is provided in line 19. `action->Get("default")` returns an `OpenGymDataContainer`. Therefore, it is necessary to dynamically cast this `OpenGymDataContainer` to the type that was sent by the `AgentApplication` (e.g. with `GetObject<OpenGymBoxContainer<int>>()`). If the content of action at key: "default" doesn't match the type passed to `GetObject`, a null pointer will be returned even if its only a mismatch in the provided data type for `OpenGymBoxContainer`.

Note: Make sure that the `OpenGymDictContainer` `action` actually contains the key passed by `action->Get("default")`. The `AgentApplication::InitiateAction` will always wrap the received action from the MARL interface into an `OpenGymDictContainer` with the key "default". However, if this method was overridden in a child class, a different key is possible.

```
1 class YourActionApp : public ActionApplication
2 {
```

(continues on next page)

(continued from previous page)

```

3 public:
4     YourActionApp(){};
5     ~YourActionApp() override{};
6
7     static TypeId GetTypeId()
8     {
9         static TypeId tid = TypeId("ns3::YourActionApp")
10            .SetParent<ActionApplication>()
11            .SetGroupName("defiance")
12            .AddConstructor<YourActionApp>();
13
14         return tid;
15     }
16
17     void ExecuteAction(uint32_t remoteAppId, Ptr<OpenGymDictContainer> action)
18     override
19     {
20         // auto m_objectActionIsPerformedOn = DynamicCast<objectActionIsPerformedOn>
21         (GetNode());
22         auto act = action->Get("default")->GetObject<OpenGymBoxContainer<int>>()->
23         GetValue(0);
24
25         m_objectActionIsPerformedOn->SetValue(acc);
26     }
27
28     void SetObservationApp(Ptr<ActionObject> object)
29     {
30         m_objectActionIsPerformedOn = object;
31     }
32
33 private:
34     Ptr<ActionObject> objectActionIsPerformedOn;
35 };

```

To perform the action, the `ActionApplication` needs a reference to the object it performs the action on. One solution would be to pass it to the application as seen in line 24-27. Alternatively, the `ActionApplication` could access the node it is installed on.

2.6 Communication between RL-Applications

2.6.1 Add interfaces

To properly use the RL applications, connect them to one another via the *ChannelInterface*. The `RLApplication` interface provides the method `RLApplication::AddInterface` to register a `ChannelInterface`. Two applications can be connected over multiple instances of `ChannelInterface`, enabling potential multipath functionality. To index the different `ChannelInterface` between two applications, an `interfaceId` has to be provided. `RLApplicationId` in combination with the `interfaceId` represents a unique identifier for a connection between two instances of `RLApplication`.

`AddInterface` also sets up necessary callbacks for receiving messages.

```

1 //code to create your agent
2 RLApplicationHelper helper(TypeId::LookupByName("ns3::YourAgentClass"));
3 helper.SetAttribute("StartTime", TimeValue(Seconds(0)));

```

(continues on next page)

(continued from previous page)

```

4 helper.SetAttribute("StopTime", TimeValue(Seconds(10)));
5 RlApplicationContainer agentApps = helper.Install(agentNode);
6
7 //code to create your observationApp
8 helper.SetTypeId("ns3::YourObservationApp");
9 RlApplicationContainer observationApps = helper.Install(obsNode);
10
11 RlApplicationId remoteIdObservationApp = DynamicCast<YourObservationApp>
12   ↳ (observationApps.Get(0))->GetId();
13 Ptr<YourAgentClass> agent = DynamicCast<YourAgentClass>(agentApps.Get(0));
14 uint interfaceAtAgentId = agent->AddInterface(remoteIdObservationApp, ↳
15   ↳ ptrToChannelInterface);
16
17 RlApplicationId remoteAgentId = DynamicCast<YourAgentClass>(agentApps.Get(0))->
18   ↳ GetId();
19 Ptr<YourObservationApp> obsApp = DynamicCast<YourObservationApp>(observationApps.
20   ↳ Get(0));
21 uint interfaceAtObservationId = obsApp->AddInterface(remoteAgentId, ↳
22   ↳ ptrToChannelInterface);

```

Note that the functionality of this method is only provided for foreseen connections of the framework. For example it is necessary that an AgentApplication can exchange data with all other types of RlApplications. Therefore the call of AgentApplication::AddInterface will succeed as long as the provided RlApplicationId::ApplicationType matches any of the following:

- OBSERVATION
- REWARD
- AGENT
- ACTION

However, if one tries to add a ChannelInterface to an ObservationApplication that is connected to another ObservationApplication, the method would result in an error because the exchange between two ObservationApplication is deliberately excluded in the design of *ns3-defiance*.

When adding the ChannelInterface, the application can derive the ApplicationType from the RlApplicationId. This allows the application to properly handle the connection.

After registering the ChannelInterface, the RlApplication is ready to send.

2.6.2 Send

Call this method to send data over a registered ChannelInterface. Note that the different RlApplications often wrap the RlApplication::Send for general use cases. Therefore, refrain from using RlApplication::Send and use the respective appropriate method offered by each application instead. These methods often ensure additional necessary prerequisites for proper communication (e.g. registering callbacks).

Even though these wrapped methods differ in their functionality they are all called in a similar manner. There are always 3 arguments: Ptr<OpenGymDictContainer> data, uint32_t appId, uint32_t interfaceIndex. The first argument is required - the data that is supposed to be sent. The second argument is the appId. If provided, the data will only be sent to the RlApplication that has a matching RlApplicationId::applicationId. The third argument the interfaceIndex can be specified alongside the RlApplicationId::applicationId. This ensures that only a specific ChannelInterface is used. The index of an interface is returned by the AddInterface method. If the interfaceIndex is not set, all interfaces between the two applications are used. Similarly, if the appId is not set the data is sent to all registered applications of that type over all interfaces.

```
// method to send actions from agent to action app
uint interfaceIdActionApp = agentApp->AddInterface(remoteActionId,
↳ptrToChannelInterface);

Ptr<OpenGymDictContainer> action = /* create DictContainer */

//send to all
SendAction(action);
//send to specific application
SendAction(action, remoteActionId);
//send to specific application over specific channelInterface
SendAction(action, remoteActionId, interfaceIdActionApp);
```

2.6.3 AgentApplication Communication

The `AgentApplication` can communicate with applications of any `RlApplicationId::ApplicationType`:

- `OBSERVATION`
- `REWARD`
- `AGENT`
- `ACTION`

See [Add interfaces](#) for more information on how to set it up.

To fulfill its functionality, the `AgentApplication` is equipped with two methods - `SendAction` and `SendToAgent`. They are invoked as described in [Send](#). `SendAction` only sends to applications of type `ACTION`, while `SendToAgent` only sends to applications of type `AGENT`.

2.6.4 RewardApplication Communication and ObservationApplication Communication

Both applications only allow communication to applications of type `AGENT`. See [Add interfaces](#) on how to add interfaces.

The interface of `RewardApplication` and `ObservationApplication` offers a `Send` method (through their parent class `DataCollectorApplication`) that works as described in [Send](#). The passed data should be used by the agent to determine the reward or update its observation.

2.6.5 ActionApplication Communication

The `ActionApplication` only allows applications of type `AGENT` to be added. See [Add interfaces](#) on how to add interfaces.

It doesn't wrap the `Send` method because it is not supposed to send, but only receive.

DATA HISTORY CONTAINER

The data history container is used for storage of data in the `AgentApplication`, specifically the latest observations received from `ObservationApplications` and latest rewards received from `RewardApplications`. When creating the history container, specify how much data it should store before deleting old data. It is possible to also specify whether the *ns-3* simulation time should be tracked with every data entry. If the usage of another history container is desired somewhere else, create a new instance of `HistoryContainer`. This can be useful for e.g. inter-agent communication.

The data container generally accepts every form of `OpenGymDictContainers`, but when the included aggregation functions like average, minimum or maximum over the last *n* entries are used, the aggregation functions will assume `OpenGymDictContainers` with `OpenGymBoxContainers` inside for them to work.

By way of the observation history container: It has an individual queue for each `ObservationApplication` that is connected to the `AgentApplication`. It also consists of a queue that contains all observations across `ObservationApplications`. The same applies to the reward history container, but with `RewardApplications`.

In order to add data to the history container, call the method `ns3::HistoryContainer::Push(ns3::Ptr<ns3::OpenGymDictContainer> data, uint id)`, which will add the data to the queue specified through `id`. This doesn't need to be done manually though, as the `AgentApplication` will automatically add the data to the history container when received from the `ObservationApplications` and `RewardApplications`. In order to do the same with agent messages, define a new history container and fill it accordingly in a method derived from `AgentApplication::OnRecvFromAgent`.

In order to get data from the history container, call the method `HistoryContainer::GetNewestByID(uint id, uint n)`, which will return the data from the queue specified through `id`. If necessary, use *n* to specify the number of entries to retrieve. If the newest data across all queues is needed, call the method `HistoryContainer::GetNewestOfCombinedHistory(uint n)`, which will return the latest *n* entries across all queues. Note that this might not retrieve evenly distributed numbers of entries from the queues, but rather the overall newest entries because different queues might be filled at different rates.

To get the average, minimum or maximum over the last *n* entries, call the method `HistoryContainer::AggregateNewest(uint id, uint n)`, which will return the average of the last *n* entries from the queue specified through `id`. This way, we can access the average, minimum or maximum of the last *n* entries for each key of the `OpenGymDictContainer`.

It makes sense to retrieve the data from the history container in the `AgentApplication` after the `AgentApplication` has received data from the `ObservationApplications` and `RewardApplications`. Thus, the methods `void OnRecvObs(uint id)` override and `void OnRecvRew(uint id)` override are the right place to retrieve the latest observations and rewards, respectively, or to do other calculations.

For example, retrieve the newest observation from the history container with ID 0 like this:

```
void OnRecvObs(uint id) override
{
    auto obs = m_obsDataStruct.GetNewestByID(0)
        ->data->Get("floatObs")
}
```

(continues on next page)

(continued from previous page)

```
                ->GetObject<OpenGymBoxContainer<float>>>();  
    m_observation = obs;  
}
```

The following code creates a mapping of `AggregatedInfo` of the last 10 entries for each key of the `OpenGymDictContainer`, providing access to average, minimum and maximum values:

```
auto agg = m_obsDataStruct.AggregateNewest(0, 10);  
auto min = agg["floatObs"].GetMin();  
auto max = agg["floatObs"].GetMax();  
auto avg = agg["floatObs"].GetAvg();
```

CHANNELINTERFACE

The channel interface is an abstraction for the communication between RL applications. It is used to send data as described in *Send*.

4.1 Overview

The channel interface is an abstracted communication channel for `RLApplications`. It provides an asynchronous, non-blocking API and uses callback mechanisms similar to the *ns-3* Socket API. It also handles serialization and deserialization of outgoing and incoming data.

The channel interface is designed to simplify communication between `RLApplications` and eliminate the overhead of creating and connecting sockets for each application. Read more about it in the Design Documentation. Additionally, the channel interface is extendable, allowing to create custom channel interfaces for other communication protocols.

The recommended way to connect and use channel interfaces is with the `CommunicationHelper` which handles the creation and connection process between `RLApplications`. Find more information in the Helper Documentation.

The channel interface sends and receives `OpenGymBoxContainer`. This makes it primarily suited for sharing observations, rewards, and actions between `RLApplications`, adhering to the OpenAI Gym API. However, due to the versatility of the `OpenGymBoxContainer`, it can be used to share arbitrary data between applications.

We provide two different channel interface implementations. Please note that **different** channel interface implementations are not interconnectable.

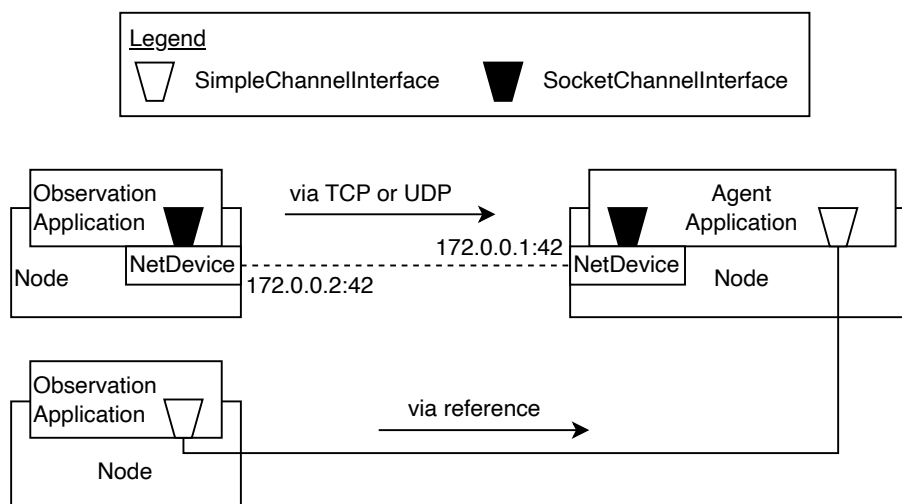


Fig. 1: Communication via `SimpleChannelInterface` and `SocketChannelInterface`

4.2 Usage

First, create two `ChannelInterface`s, one for each `RLApplication` that will communicate with each other. Connect the two `ChannelInterface` objects using the `ChannelInterface::Connect` method. Afterwards, send data to the remote `RLApplication` using `ChannelInterface::Send`.

To handle received data, add a callback function to the channel interface with the `ChannelInterface::ConnectAddRecvCallback` method. This callback function will be called when new data arrives, with the deserialized data as an `OpenGymBoxContainer`.

If a callback function is no longer needed, remove it using the `ChannelInterface::ConnectRemoveRecvCallback` method. Add as many callback functions as needed. They will be called in the order they were added.

Disconnect the two `ChannelInterface` objects with the `ChannelInterface::Disconnect` method. For that, provide the specific callback function that shall be remove.

Check the connection status of the channel interface using `ChannelInterface::GetConnectionStatus`. It returns an element of the following enum:

```
enum ConnectionStatus
{
    DISCONNECTED,
    CONNECTING,
    CONNECTED,
};
```

4.3 SimpleChannelInterface

The `SimpleChannelInterface` simulates communication between `RLApplications` without using the underlying network simulation. It is primarily intended for debugging or simulating communication without the overhead of a full network simulation. It does not provide a realistic simulation of network communication and should not be used for performance evaluation. However, set a network delay to simulate network latency if needed.

Here is an example of how to use the `SimpleChannelInterface`:

```
// the simple interface does not need any configuration or parameters
auto interfaceSimple0 = CreateObject<SimpleChannelInterface>();
auto interfaceSimple1 = CreateObject<SimpleChannelInterface>();

// create a callback function which prints the contents of the OpenGymDictContainer
auto recvCallback = Callback<void, Ptr<OpenGymDictContainer>>([
    ](Ptr<OpenGymDictContainer> msg) { NS_LOG_INFO(msg->Get("box")); });

// add the callback function to the channel interfaces, both should just print the
// received data
interfaceSimple0->AddRecvCallback(recvCallback);
interfaceSimple1->AddRecvCallback(recvCallback);

// add a simple network delay of 0.1 seconds
interfaceSimple0->SetPropagationDelay(Seconds(0.1));
interfaceSimple1->SetPropagationDelay(Seconds(0.1));

// connect the two channel interfaces with each other in the simulation after 0.1
// seconds
```

(continues on next page)

(continued from previous page)

```

Simulator::Schedule(Seconds(0.1),
                    &SimpleChannelInterface::Connect,
                    interfaceSimple0,
                    interfaceSimple1);

/* helper method to creates a OpenGymDictContainer
   with a OpenGymBoxContainer named "box" with a float value */
Ptr<OpenGymDictContainer>
CreateTestMessage(float value)
{
    Ptr<OpenGymDictContainer> msg = Create<OpenGymDictContainer>();
    Ptr<OpenGymBoxContainer<float>> box = Create<OpenGymBoxContainer<float>>();
    box->AddValue(value);
    msg->Add("box", box);
    return msg;
}

// send the OpenGymDictContainer from interfaceSimple0 to interfaceSimple1 after 1
↪second */
Simulator::Schedule(Seconds(1),
                    &SimpleChannelInterface::Send,
                    interfaceSimple0,
                    CreateTestMessage(0));

```

This example creates two `SimpleChannelInterface` objects and connects them. After 1 second, it sends a message from one interface to the other. Due to the 0.1 second network delay, the message is printed by the receiving interface after 1.1 seconds.

4.4 SocketChannelInterface

The `SocketChannelInterface` uses sockets to communicate between `RLApplications`. It utilizes *ns-3* sockets under the hood and is the recommended way to simulate realistic network communication.

The network scenario and topology should ensure that the `RLApplications` can communicate with each other, for example, via the Internet or a local network. The channel interface itself does not handle the network communication; it only provides the API for communication.

If other communication methods are required, create a custom channel interface and implement it accordingly.

Here is an example of how to use the `SocketChannelInterface`:

```

// create nodes
NodeContainer nodes;
nodes.Create(2);

// create a point-to-point helper
PointToPointHelper p2p;
p2p.SetDeviceAttribute("DataRate", StringValue("5Mbps"));
p2p.SetChannelAttribute("Delay", StringValue("2ms"));

// create devices and install them on nodes
NetDeviceContainer devices;
devices.Add(p2p.Install(nodes.Get(0), nodes.Get(1)));

// assign IP addresses

```

(continues on next page)

(continued from previous page)

```

InternetStackHelper internet;
internet.Install(nodes);

Ipv4AddressHelper address;
address.SetBase("10.1.1.0", "255.255.255.0");
Ipv4InterfaceContainer interfaces = address.Assign(devices);

// get the protocol type id for the protocol we want to use
auto tcpProtocol = TcpSocketFactory::GetTypeId();

// create the interfaces with the nodes and the protocol
Ptr<SocketChannelInterface> interfaceTcp0_1 =
    CreateObject<SocketChannelInterface>(nodes.Get(0), interfaces.GetAddress(0),
    ↪tcpProtocol);
Ptr<SocketChannelInterface> interfaceTcp1_0 =
    CreateObject<SocketChannelInterface>(nodes.Get(1), interfaces.GetAddress(1),
    ↪tcpProtocol);

// create a callback function which prints the contents of the OpenGymDictContainer
auto recvCallback = Callback<void, Ptr<OpenGymDictContainer>>({
    [](Ptr<OpenGymDictContainer> msg) { NS_LOG_INFO(msg->Get("box")); } });

// add the callback function to the channel interfaces, both should just print the
    ↪received data
interfaceTcp0_1->AddRecvCallback(recvCallback);
interfaceTcp1_0->AddRecvCallback(recvCallback);

// connect the two channel interfaces with each other in the simulation after 0.1
    ↪seconds
Simulator::Schedule(Seconds(0.1),
                    &SocketChannelInterface::Connect,
                    interfaceTcp0_1A,
                    interfaceTcp1_0);

/* helper method to creates a OpenGymDictContainer
   with a OpenGymBoxContainer named "box" with a float value */
Ptr<OpenGymDictContainer>
CreateTestMessage(float value)
{
    Ptr<OpenGymDictContainer> msg = Create<OpenGymDictContainer>();
    Ptr<OpenGymBoxContainer<float>> box = Create<OpenGymBoxContainer<float>>();
    box->AddValue(value);
    msg->Add("box", box);
    return msg;
}

// send the OpenGymDictContainer from interfaceTcp0_1 to interfaceTcp1_0 after 1
    ↪seconds */
Simulator::Schedule(Seconds(1),
                    &SocketChannelInterface::Send,
                    interfaceUdp0_1,
                    CreateTestMessage(1));

```

This example creates two `SocketChannelInterface` and connects them. After 1 second, it sends a message from one interface to the other and prints the received message after approximately 1.02 seconds (because of the 20ms network delay).

4.5 Custom Channel Interface

If necessary, implement and use a custom channel interface to use alternative communication protocols or methods for communication between `RLApplications`.

To create a custom channel interface, inherit from the abstract base class `ChannelInterface` and implement its corresponding methods.

HELPER CLASSES

5.1 RlApplicationHelper

As previously introduced, the DEFIANCE framework is mainly structured around user specific `RlApplications`. They are derived from their specific base classes (e.g. `AgentApplication`) and communicate relevant information with one another during the simulation.

To simplify the creation of their instances, the `RlApplicationHelper` class is provided. As with the typical helper classes already present in *ns-3*, it makes the creation of the applications more intuitive.

The following example demonstrates how the `RlApplicationHelper` can be used.

```
RlApplicationHelper helper (TypeId::LookupByName ("ns3::MyObservationApp"));

// the helper allows to set attributes for the applications
// this is persistent for all the applications that will be created afterwards
helper.SetAttribute("StartTime", TimeValue(Seconds(0)));
helper.SetAttribute("StopTime", TimeValue(Seconds(10)));

RlApplicationContainer observationApps = helper.Install(myNodes1);

helper.SetTypeId("ns3::MyRewardApp");
RlApplicationContainer rewardApps = helper.Install(myNodes2);

helper.SetTypeId("ns3::MyActionApp");
RlApplicationContainer actionApps = helper.Install(myNodes3);

helper.SetTypeId("ns3::MyAgentApp");
RlApplicationContainer agentApps = helper.Install(myNodes4);
```

This example shows the main features of the `RlApplicationHelper`. First of all, it wraps the created application instances in an `RlApplicationContainer`. This container can be used like the standard *ns-3* `ApplicationContainer` to access or iterate over the applications but does not require to cast the applications each time that DEFIANCE-specific functionality is required. Secondly, the helper allows to set attributes for the applications. This enables work with the `TypeId` system, which makes it easy to set default arguments and to work with command line arguments. In the example above, the helper is used to create different types of applications but sets the same start and stop time for all of them.

Note: The `RlApplicationHelper` is not limited to the applications that are provided by the DEFIANCE framework. It can be used with any application that is derived from the `RlApplication` class.

5.2 CommunicationHelper

The natural extension to the `RlApplicationHelper` is the `CommunicationHelper`. It can work with `RlApplicationContainers` to create communication channels between the applications and configure them accordingly. The `CommunicationHelper` simplifies this procedure and reduces the risks of bugs.

First, create an instance of `CommunicationHelper` and set the different applications:

```
1 CommunicationHelper commHelper = CommunicationHelper();
2
3 commHelper.SetObservationApps(observationApps);
4 commHelper.SetAgentApps(agentApps);
5 commHelper.SetRewardApps(rewardApps);
6 commHelper.SetActionApps(actionApps);
7 commHelper.SetIds();
```

The different `Set` methods expect an object of type `RlApplicationContainer`. See chapter `ApplicationHelper` for more information on how to create one. After the helper received all `RlApplicationContainers`, the IDs of these applications need to be assigned (line 7). The IDs are used to identify the instances of `RlApplication` and are required for the next step.

Once the IDs are assigned, the actual connection can be configured. This can be done by passing a vector of type `CommunicationPair` to the `CommunicationHelper`. To create an instance of `CommunicationPair`, the IDs of the two `RlApplications` and a `CommunicationAttributes` object have to be provided. The `CommunicationAttributes` object describes the type of connection. If no argument is passed, a `SimpleChannelInterface` is created. To create a socket connection via TCP or UDP, a `SocketCommunicationAttributes` object with `TypeId` protocol set accordingly can be passed as `CommunicationAttributes`. The following code is a simple example that creates `CommunicationPairs` of different types.

```
1 // UDP
2 CommunicationPair actionCommPair = {
3     actionApps.GetId(0),
4     agentApps.GetId(0),
5     SocketCommunicationAttributes{"7.0.0.2", "1.0.0.2", UdpSocketFactory::GetTypeId()}
6     ↪};
7
8 //TCP
9 CommunicationPair observationCommPair = {
10     observationApps.GetId(0),
11     agentApps.GetId(0),
12     SocketCommunicationAttributes{"7.0.0.2", "1.0.0.2", ↪
13     ↪TcpSocketFactory::GetTypeId()}};
14
15 //SIMPLE
16 CommunicationPair actionCommPair = {actionApps.GetId(0),
17                                     agentApps.GetId(0),
18                                     {}};
```

The method `GetId(i)` allows to retrieve the `RlApplicationId` by passing the index `i` to the `RlApplicationContainer` (as used in e.g. line 3–4). When creating `SocketCommunicationAttributes`, the passed IP addresses have to match the addresses of the node the application is installed on.

Once these `CommunicationPairs` are created, collect them in a vector and pass it to `CommunicationHelper::AddCommunication` as a parameter. Finally, the configuration can be finished by calling `Configure` on the `CommunicationHelper`. Now all channel interfaces are created accordingly, ready for sending and receiving data. An explanation of the `Configure` method can be found in section *Helper* of the design documentation.